

ParaGnosis Tool Demo

ParaGnosis is a C++ weighted model counting toolset for Linux. Its implementation is based on [\[1,2,3,4\]](#). We have also added a significant number of Bayesian networks to play with (under `./data/net`). This demo will first introduce the user to the toolset and the input formats. Then we will demonstrate three ways to use the tool:

- Compile a Bayesian network to a knowledge base in different target languages.
- Compare the results of the compilation process for different target languages.
- Visualize the compilation result using dot
- Perform various inference queries on the compiled knowledge base:
 - marginalization,
 - conditional probabilities, and
 - posteriors.
- Compare marginalization results on different target languages.

The tool consists of the following command-line tools:

- `bn-to-cnf` : a c++ tool to create Conjunctive Normal Form (CNF) encodings from a Bayesian network.
- `bnc` : a c/c++ **B**ayesian **N**etwork **C**ompiler for multiple target representations.
- `bnmc` : a c++ **B**ayesian **N**etwork **M**odel **C**ounter.
- `pg` : a **ParaG**nosis user friendly interface to the tools above, written in Python.

The currently supported target languages are:

- Weighted Positive Binary Decision Diagrams (WPBDD)
- Weighted Positive Multi-Valued Decision Diagrams (WPMDD)
- Tree-driven Weighted Positive Multi-valued Decision Diagrams (TD-WPMDD)

The demo

The `pg` script is installed system wide, and directly available from the command-line by opening a terminal. A shortcut to open a terminal is `ctrl-alt-t`.

Encoding

Show a list of available Bayesian networks

The toolset comes with a comprehensive list of Bayesian networks to play with. To get a list of available networks, type:

```
> pg --list

3nt
4sp
6hj
6nt
aggregate
alarm
...
```

Any of the shown names can be used as input for the `pg` script. You can also provide a locally stored Bayesian network filename with `.net` extension (HUGIN format).

Show encoding statistics for the *asia* network

```
> pg encode asia

...

Variables      : 8
Probabilities  : 36
Deterministic  : 8
Unsatisfiable  : 4
Literals       : 16
Clauses        : 52
Literal/clauses : 2.23
Clause sizes   : 1-3
```

Compiling a network

Compile *asia* to a TD-WPMDD

```
> pg compile asia

...

FINAL RESULT:

Spanning tree      : 0.011ms
Compilation        : 0.015ms
Total Or #nodes    : 23
Total And #nodes   : 6
Total time         : 0.000s
Total time         : 0.026ms

Total #nodes       : 29
Total #edges       : 58
Total #operators   : 144
```

Compile *asia* to a WPBDD

```
> pg compile asia --method wpbdd

...

FINAL RESULT:

Compiled CPTs in   : 0.000s
Conjoined CPTs in  : 0.000s
Total time         : 0.000s
Total time         : 0.244ms

Total #nodes       : 45
Total #edges       : 90
Total #operators   : 124
```


Compare compilation between WPBDD, a WPMDD, and a TD-WPMDD

```
> pg compile asia --method wpbdd mg tdmg
```

...

nr	type	seconds	milliseconds	speed-down	operators	nodes	edges
0	TDMG	0.000	0.031	1.000	144	29	58
1	MG	0.000	0.089	2.871	132	22	44
2	WPBDD	0.000	0.155	5.000	124	45	90

The following is the *icy roads* Bayesian network that we want to compile



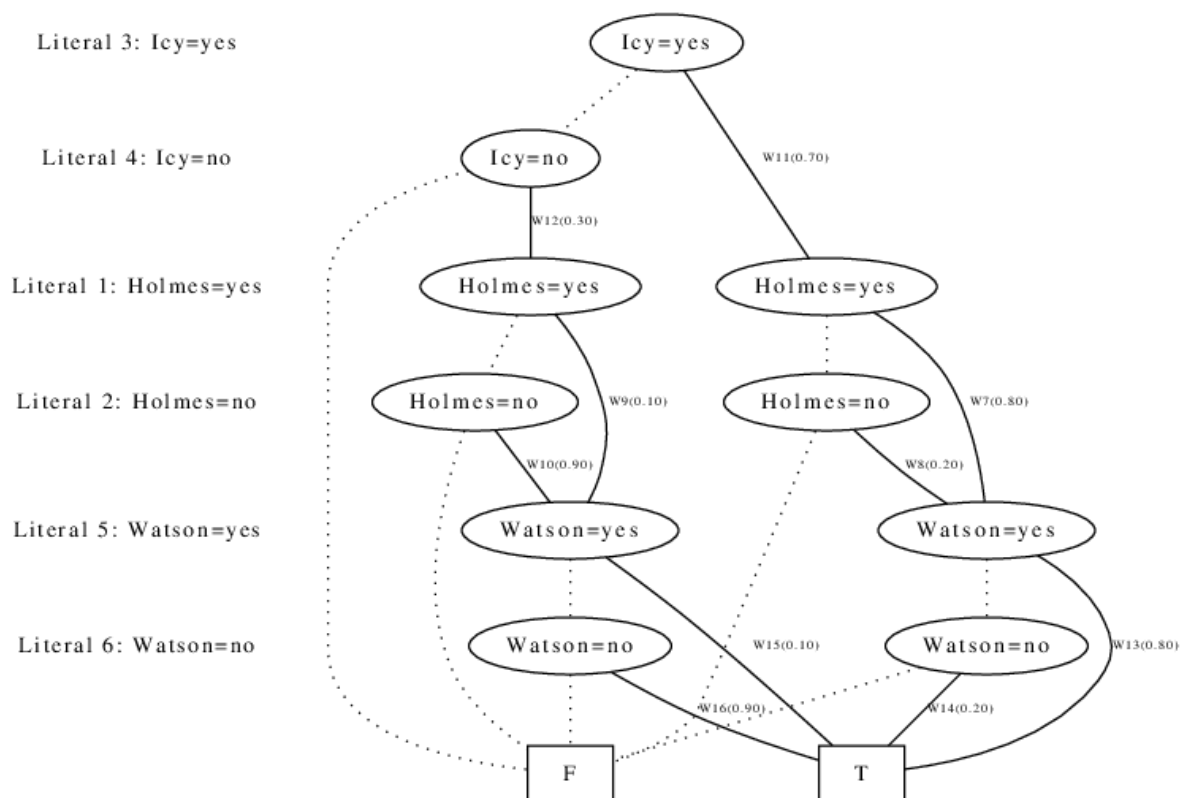
$P(X_{Icy})$:
 yes 0.7
 no 0.3

$P(X_{Holmes}|X_{Icy})$:
 yes no
 yes 0.8 0.2
 No 0.1 0.9

$P(X_{Watson}|X_{Icy})$:
 yes no
 yes 0.8 0.2
 no 0.1 0.9

Directly visualize the WPBDD for the *icy_roads* network


```
> pg compile icy_roads --method wpbdd --dot
```



Directly visualize the TD-WPMDD

```
> pg compile icy_roads --method tdmg --dot
```


Ordering



Directly visualize the TD-WPMDD

```
> pg compile icy_roads --method tdmg --dot
```

Ordering



Perform Inference

Verify *icy_roads* posteriors, previously shown

```
> pg inference --posteriors="" icy_roads

...

Holmes=yes: 0.590000
Holmes=no: 0.410000
Icy=yes: 0.700000
Icy=no: 0.300000
Watson=yes: 0.590000
Watson=no: 0.410000
```

Run every possible marginalization on a network using TD-WPMDD (press ctrl-c to stop)

```
> pg inference asia

...

  nr |           type |  cores |      queries | milliseconds |      ms/q
----|-----|-----|-----|-----|-----
  0 | TDMULTIGRAPH |      1 |      18360 |      8.7516 |      0.0005
```

Compare WPBDD inference speed with TD-WPBDD

```
> pg inference asia --method wpbdd tdmg --compare

...

  nr |           type |  cores |      queries | milliseconds |      ms/q | speed-down
----|-----|-----|-----|-----|-----|-----
  0 | WPBDD         |      1 |      18360 |      5.7327 |      0.0003 |      1.00
  1 | TDMULTIGRAPH |      1 |      18360 |     10.5697 |      0.0006 |      1.84
```

Compute conditional probability $P(\text{tub} \mid \text{bronc} = \text{yes}, \text{smoke} = \text{yes})$ for the *asia* network

```
> pg inference asia --evidence='bronc=yes,smoke=yes' --posteriors='tub'

...

tub=yes: 0.010400
tub=no: 0.989600
```

Compute posteriors of `lung` and `xray` for evidence `bronc = yes`, and `smoke = yes`

```
> pg inference asia --evidence='bronc=yes,smoke=yes' --posteriors='lung,xray'

...

lung=yes: 0.100000
lung=no: 0.900000
xray=yes: 0.151705
xray=no: 0.848295
```

Compute posteriors all non-observed variables, for evidence `bronc = yes`, and `smoke = yes`.

```
> pg inference asia --evidence='bronc=yes,smoke=yes'

...

asia=yes: 0.010000
asia=no: 0.990000
dysp=yes: 0.810936
dysp=no: 0.189064
either=yes: 0.109360
either=no: 0.890640
lung=yes: 0.100000
lung=no: 0.900000
tub=yes: 0.010400
tub=no: 0.989600
xray=yes: 0.151705
xray=no: 0.848295
```

References

- [1] G.H. Dal, A.W. Laarman, A. Hommerso and P.J.F. Lucas, "A *Compositional Approach to Probabilistic Knowledge Compilation*", in International Journal of Approximate Reasoning, vol 138:38-66, 2021.
- [2] G.H. Dal, A.W. Laarman and P.J.F. Lucas, "*Parallel Probabilistic Inference by Weighted Model Counting*", in Proceeding of the International Conference on Probabilistic Graphical Models, PMLR, vol 72:97-108, 2018.
- [3] G.H. Dal, S. Michels and P.J.F. Lucas, "*Reducing the Cost of Probabilistic Knowledge Compilation*", in Proceedings of Machine Learning Research, volume 73, pages 41-152, 2017.
- [4] G.H. Dal and P.J.F. Lucas, "*Weighted Positive Binary Decision Diagrams for Exact Probabilistic Inference*", in Journal of Approximate Reasoning, volume 90, pages 411-432, 2017.